Supplementary Material

Individual aDDM Fitting.

The individual-level fitting procedure was very similar to the group-level fitting (see Methods), with a few key differences. First, we split the data into 5 bins in the food-risk, money-risk and social tasks at the individual level. Second, in the first iteration of fits, we tested 20,000 parameter combinations for each subject in each task with 100 simulations per trial, with wide initial ranges for the parameters as follows: d = [0,0.001], σ = [0,0.05], θ = [0,1], t_{er} = [0,3000]. In all subsequent iterations, the range for θ was set to [0,1].

We could not fit some subjects (four in two-food, three in food-risk, two in money-risk, and four in social) with such a fine division of trials (dwell-time terciles and five utility difference bins). If a subject could not be fit with that many divisions, we re-attempted their fits using a median split on dwell time. We also re-attempted their fits using a coarser split on utility difference (three bins instead of five). Most of the problematic subjects were remedied by one of these two revisions, but a few subjects (one in the two-food task and one in the social task) required an even coarser split. For the subject in the food task, we used two utility difference bins (i.e. high and low), plus a median split in dwell time advantage. The subject in the social task could not be fit with this minimal split, so they were excluded from future fitting attempts.

After establishing the necessary fitting technique for each subject, we ran three more iterations. We identified the range of parameters combinations in the top 10%, according to MLE, and used these values to set the bounds for the second iteration (for d, σ and t_{er}). In this second iteration, we ran 5000 parameter combinations for each subject in each task. The third and fourth iterations used parameter ranges from the top 1% of previously-attempted parameter sets and comprised 2500 parameter combinations each. Finally, we identified the best fitting parameters (out of 30,000 potential combinations) for each subject. Many subjects were fit quite well by this procedure, while others were clearly not (Fig. S4). It is unclear whether this is a shortcoming of the aDDM, the utility functions, or the fitting procedure. More work is needed to fully understand these issues.

Theta vs. Beta

A question of interest in this study is the relationship between the attentional discounting parameters across tasks. As seen in the previous modeling results, there are clear behavioral and attentional consistencies across the different domains. Any aDDM consistencies in attentional discounting (θ) parameters, therefore, would add further evidence for a domain-general decision-making process, moderated by attention to the different alternatives. However, we encountered issues while fitting the model to some individual subjects, which interfered with these analyses. It is unclear why the model fitting failed for some subjects, but it could be due to choices that do not conform to the utility function, noisy behavior, issues with the goodness-of-fit metric for these subjects, or simply a model that does not work for every individual. Understanding why the modeling fails for some individuals is an ongoing topic of study. All that being said, we can see that the model provides a good fit for most of our subjects on most of the tasks (Fig. S4).

In light of these model-fitting complications, we pursued an alternative strategy to study the consistency of the attentional effect on choice. In particular, we focused on the aDDM's core feature, the effect of dwell time on choice probability, using subject-level logistic regressions. To establish the connection between the aDDM's attentional parameter θ and the logistic β coefficients in a model of choice on dwell time advantage (with *value difference* as a covariate), we simulated 101 datasets with θ values = (0, 0.01, 0.02... 1). To maintain robustness, the other parameters were randomly selected from uniform distributions with the following ranges: d =

[0.0001,0.0003], σ = [0.01,0.03], t_{er} = [100,500]. After simulating, we fit the following regression:

$$P(ChooseLeft) = \beta_0 + \beta_1(U_L - U_R) + \beta_2(D_L - D_R)$$

Then, we compared the fitted regression coefficients to the thetas that generated the simulations. There is a very strong negative correlation between θ and β, $r(99) = -0.72$, $p < 10^{-16}$. This suggests that the logistic regression coefficients are a suitable proxy for the attentional influence observed in choices at the individual level.

In fact, when we look at this correlation in each of our tasks, the results are encouraging. We find that the dwell-time coefficients are mostly correlated with the aDDM thetas: two-food: $r(42) = -0.53$, $p = 0.0003$; food-risk: $r(42) = -0.47$, $p = 0.001$; money-risk: $r(34) = -0.42$, $p = 0.01$, social (combined): $r(34) = -0.29$, $p = 0.08$; social (pro-social): $r(20) = -0.28$, $p = 0.20$; social (selfish): $r(12) = -0.33$, $p = 0.25$.

Alternative aDDM Simulations

The aDDM simulations included in the main text reflect the actual dwell patterns of subjects, while the simulations in Figs. S5-S8 follow the method of Krajbich et al. (2010) and use randomly sampled dwell times and first dwell locations from the task-level data across subjects and value differences.

Transitions and Dwells

Subjects varied in their transition patterns, within and across tasks (as in Fig. 8 in the main text). On the whole, subjects tended to make more within-option transitions in the food-risk and money-risk tasks, while they made more between-option transitions in the social task (especially the selfish subjects). There was also ample variability in the number of dwells during a given choice at the trial, subject, and task levels, as well as in the location of the subjects' first dwells (Figs. S9-S11).

Final Dwell Analysis

The location of the final dwell is also a significant predictor of choice. When we fit an individual logistic regression of choice as a function of dwell time advantage and final dwell location, we see that there are significant relationships across the tasks in the magnitude of the final dwell coefficient (Table 3 in the main text). We also estimated an individual-level model with both final dwell location and dwell time advantage included as covariates. The coefficient correlations from this more complicated model are included below. However, because of collinearity issues (i.e. the last seen option also tends to have been looked at longer), we are hesitant to make any strong claims from this latter analysis. Considering both the dwell-time and final-dwell correlations together (see Tables S1 & S2), we see significant across-task correlations at least once for every pair of tasks (9/12 significant correlations), looking at all participants. The split-group correlations are generally consistent with what we reported in the paper, though in some cases they are only present for one measure and not the other. In some cases, we do see tradeoffs between dwell-time and final-dwell effects (one large and one small), consistent with our concerns about collinearity.

Additionally, we conducted analyses of the same combined (final dwell and dwell time advantage) model described above at the group level. When we compare these results to those in the main text, we find that they do not change substantially. Specifically, an ANOVA of the FinalDwellLeft coefficients did not reveal any significant differences across the tasks, *F*(3, 74.98) = 1.21, $p = 0.31$. Additionally, all pairwise t-tests were insignificant (ps > 0.35), even without correcting for multiple comparisons (ps > 0.05). When we run a group-level mixed effects model (with a Bayesian approach), we find that the 95% confidence intervals for the FinalDwellLeft coefficient do, indeed, intersect: two-food: [1.38,2.07]; food-risk: [1.35,2.15]; money-risk: [1.66,2.48]; social (pro-social): [1.37,2.47]; social (selfish): [-0.06,2.64]. Ultimately, when we estimate the effects of dwell time and final dwell in a single model, some – but definitely not all – of the variance in choice behavior is shared by the two predictors.

Furthermore, we examined the proportion of choices to the last-seen option as a function of trial number. There were no systematic trends for subjects to make more/fewer choices for the last-seen option over time (Fig. S19; Table S4), with the exception of the selfish subjects in the social task.

Psychophysical Task

In the main text, we chose to report the proportion of trials where the subjects chose the more looked-at option because it was a model-free way to look at the effects of attention across all tasks in one test. However, we have included additional information about this psychophysical task and its relationship to fitted parameters here (Figs. S12-S14). We would expect a positive correlation in Figs. S13-S14 because a larger sharpness gradient implies a narrower scope of attention and thus, a larger influence of dwell time on choice.

For the correlations between the sharpness gradient and the beta on dwell time, we find: two-food: $r(30) = 0.41$, $p = 0.02$; food-risk: $r(30) = -0.04$, $p = 0.85$; money-risk: $r(25) = 0.56$, $p = 0.0006$; social (combined): $r(25) = 0.40$, $p = 0.04$; social (pro-social): $r(17) = 0.43$, $p = 0.07$; social (selfish): $r(6) = 0.43$, $p = 0.30$. For the correlations between the sharpness gradient and the beta on FinalDwellLeft, however, we don't see the same relationship. We find: two-food: $r(30) = 0.03$, $p = 0.88$; food-risk: $r(30) = -0.32$, $p = 0.08$; money-risk: $r(25) = 0.30$, $p = 0.13$; social (combined): $r(25) = -0.49$, $p = 0.01$; social (pro-social): $r(17) = -0.04$, $p = 0.87$; social (selfish): $r(6) = -0.21$, $p = 0.62$.

If we instead use the model with both dwell time and final dwell (which we have argued is problematic), we find a similar pattern of results. For the (Spearman) correlations between the sharpness gradient and the beta on dwell time (estimated jointly with the final dwell), we find: two-food: $r(30) = 0.32$, $p = 0.08$; food-risk: $r(30) = 0.07$, $p = 0.71$; money-risk: $r(25) = 0.57$, $p = 0.002$; social (combined): $r(25) = 0.35$, $p = 0.07$; social (pro-social): $r(17) = 0.27$, $p = 0.27$; social (selfish): $r(6) = 0.55$, $p = 0.17$. For the correlations between the sharpness gradient and the beta on FinalDwellLeft (estimated jointly with the dwell time), however, we don't see the same relationship. We find: two-food: $r(30) = 0.003$, $p = 0.98$; food-risk: $r(30) = -0.38$, $p = 0.03$; money-risk: $r(25) = 0.08$, $p = 0.70$; social (combined): $r(25) = -0.09$, $p = 0.64$; social (pro-social): $r(17) = -0.16$, $p = 0.51$; social (selfish): $r(6) = 0.31$, $p = 0.46$.

Order Effects

As stated in the main text methods, we randomized the order of the tasks across subjects. We did not observe any systematic effects of task location (i.e. whether a subject completed a given task 1^{st} /2^{nd}/3^{rd}/4^{th}), but we did observe a decrease in RT as subjects gained more experience within a task (regression of RT on trial number with clustered standard errors: two-food: $\beta = -0.0013$, $p = 0.002$; food-risk: $\beta = -0.00011$, $p = 0.95$; money-risk: $\beta = -0.0085$, $p = 10^{-9}$; social (pro-social): $\beta = -0.0037$, $p = 0.0003$; social (selfish): $\beta = -0.0036$, $p = 10^{-5}$). Accuracy declined slightly over time in the social task, but increased in the money-risk task (logistic regression of accuracy on trial number with clustered standard errors: two-food: $\beta = -0.00010$, $p = 0.32$; food-risk: $\beta = 0.000027$, $p = 0.86$; money-risk: $\beta = 0.00063$, $p = 0.001$; social (pro-social): $\beta = -0.00022$, $p = 0.05$; social (selfish): $\beta = -0.00039$, $p = 0.001$). Summaries of these analyses are included in Fig. S15 and Table S3.

Food Rating Distributions

There was ample variability in the ratings given by subjects, including subjects' average food ratings, and the standard deviation of their ratings (Fig. S16).

Additional Social Task Findings

One subject chose the pro-social option in 97.5% of trials. It is possible that this subject misunderstood the directions and mixed up the colors, intending to always choose the selfish option instead. If we exclude this subject (or instead re-code the subject as selfish), the results do not change qualitatively.

The behavioral discrepancies observed between the two groups of subjects extend to the eye-tracking results, as well. For example, selfish subjects looked only at their own amounts (the red boxes) on 66.1% (SD = 20.3%) of trials, while the mean percentage for pro-social subjects was only 13.2% (SD = 12.9%) ($t(19.706) = 8.7133$, $p = 10^{-7}$). Selfish subjects also attended to only one out of the four ROIs on 43.1% (SD = 13.4%) of trials while pro-social subjects attended to one ROI on only 12.2% (SD = 10.7%) of trials ($t(23.239) = 7.2839$, $p = 10^{-6}$).

Heuristics

There are a couple of characteristics of heuristics (in general) that we can look for in our data. Heuristics use a consistent, simple rule to make decisions, and so the choice process should look fairly similar across trials. In particular, any given heuristic should yield decisions that take roughly the same amount of time, irrespective of the subjective value difference (SVD). For the same reason, heuristics should exhibit a consistent pattern of attention and little-to-no effects of attention on choice.

Therefore, if subjects were indeed using a heuristic or shortcut in the later trials, then we would expect to find a couple of trends in the later trials. First, we would not expect their response times (RT) to depend as much on SVD. Second, we would not expect choices to depend as much on dwell-time differences.

As we highlight in the main text, the closest thing to a heuristic in our experiment is being purely selfish and thus needing to simply identify the bigger self-payoff. Therefore, below we contrast the purely selfish choices with the other tasks.

With regard to RTs, we generally find no evidence for a decreasing relationship between RT and SVD across trials (Fig. S17). The lone exception is the purely-selfish set of subjects; here the absolute effect of SVD on RT gets closer to zero as trial number increases.

There is still a significant relationship between SVD and RT in the selfish subjects, but this is certainly due to perceptual similarity effects, commonly observed in perceptual decision-making tasks. That is, when the red boxes (i.e. self-payoff amounts) are close in size, then subjects take some time to find the bigger one. As we move away from equal size boxes, RTs decrease and asymptote quickly, relative to the other tasks. At the maximum SVDs, selfish subjects' choices are more than a full second faster than the other four-stimulus tasks (Fig. 4E).

Another observation is that the money-risk task initially shows little correlation between SVD and RT, but this relationship is present in the second half of trials. However, this goes in the opposite direction predicted by the adoption of a heuristic. Moreover, as we will discuss in the next paragraph, there are no corresponding changes in the relationship between dwell-time and choice in the money-risk task.

With regard to the link between attention and choice, we again find no evidence for a decreasing relationship between dwell-time difference and choice across trials (Fig. S18). Again, the lone exception is the purely-selfish set of subjects. Unlike the other tasks, where the effect of dwell-time difference on choice is fairly constant (in the range of 1-2), in the purely selfish subjects the effect of dwell time advantage on choice decreases to 0 within ~60 trials.

We can also examine the proportion of choices for the last-seen option as a function of trial (Fig. S19). In short, there do not seem to be any strong differences across the trials. For each task, we also regressed the proportion of choices for the last-seen option on the trial number, and the only task that had a significant relationship was the selfish subjects in the social task. The results from this regression are included below (Table S4).

If subjects were adopting different heuristics, then we might also expect larger standard errors on the coefficients as trial number increases. However, we do not see any evidence for this in any set of plots.

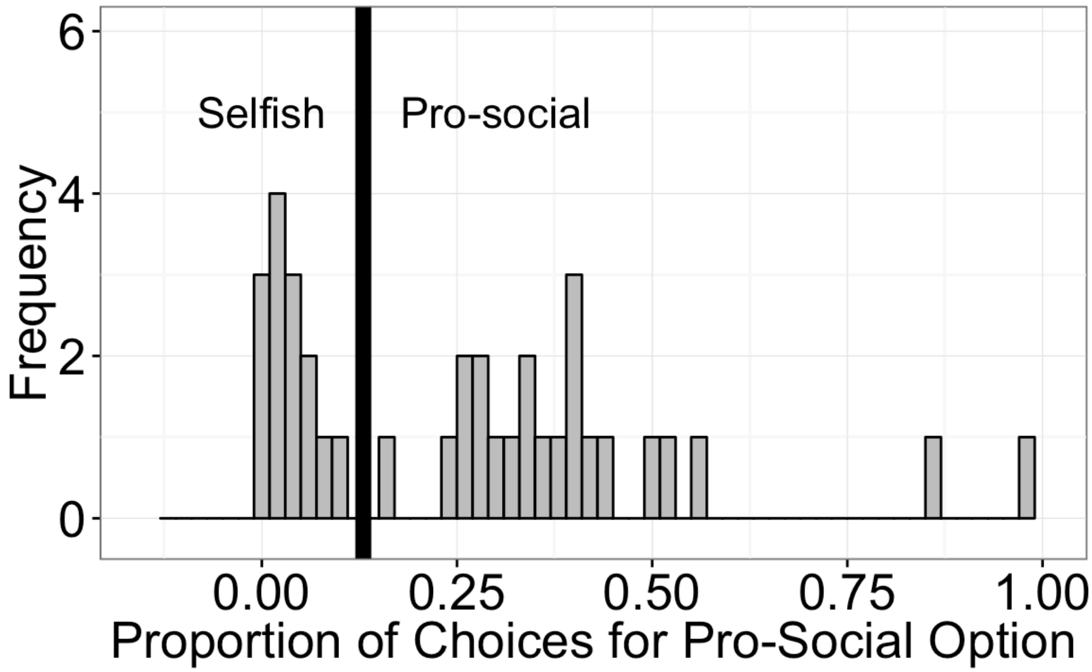

Figure S1. Proportion of pro-social choices. Histogram of choices made by subjects in the social task, with vertical bar indicating the separation into pro-social and selfish groups.

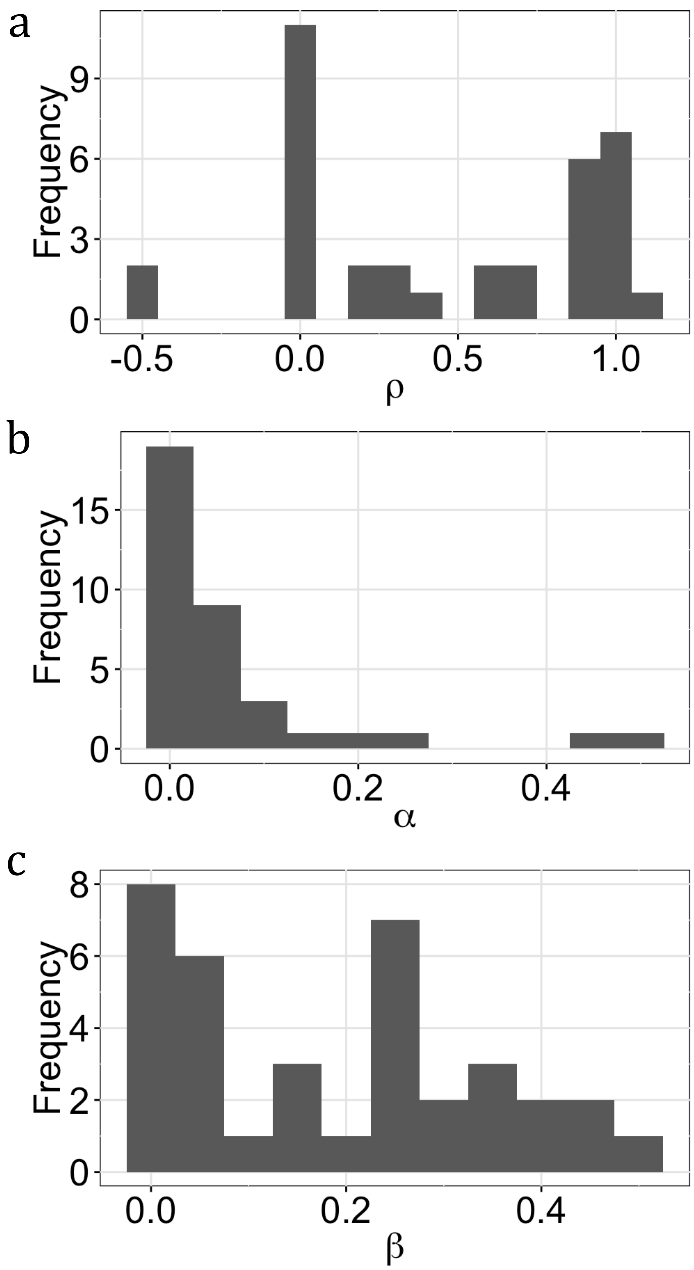

Figure S2. Utility model parameters. Histograms for best-fitting parameters in the exponential utility (a) and Charness-Rabin (b,c) models. (a) The exponential utility parameter: rho. Lower values of rho indicate greater risk aversion. (b,c) The Charness-Rabin disadvantageous and advantageous inequality parameters: alpha and beta. Higher values of each parameter indicate a greater influence of inequality on an individual's utility.

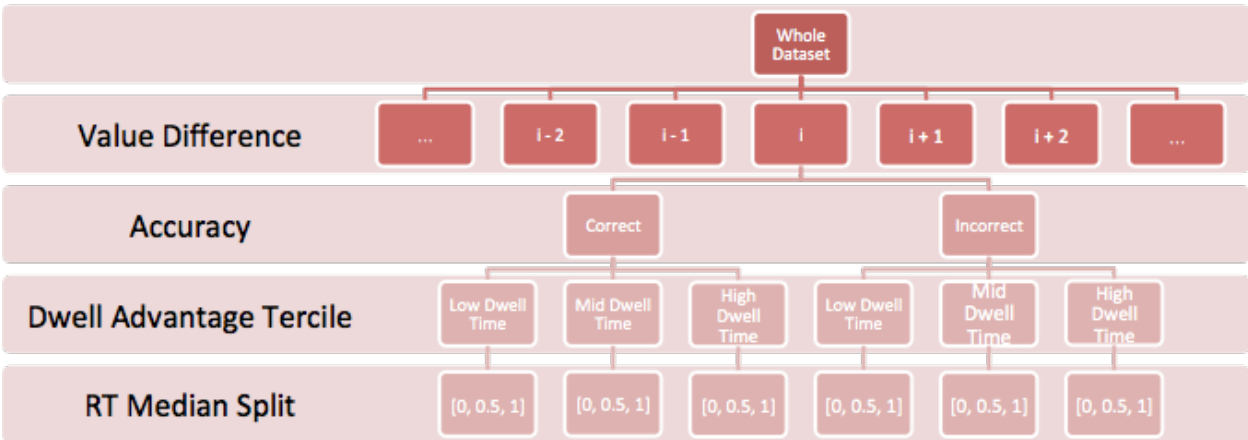

Figure S3. Binning procedure for the aDDM. Trials were split according to value difference, accuracy, dwell time advantage, and RT, except for value difference = 0, which was not split by accuracy.

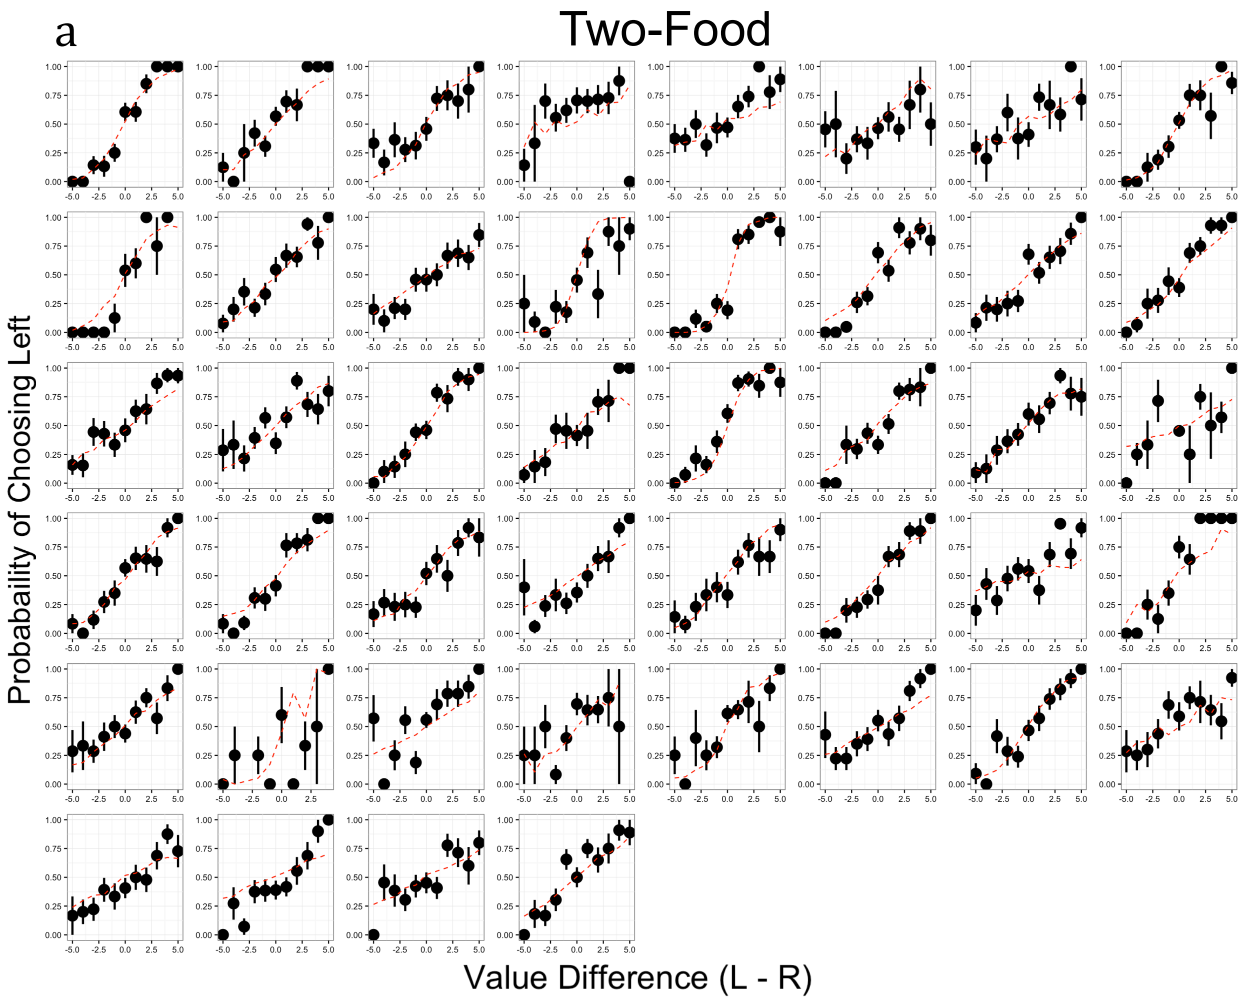
a
Two-Food
Probability of Choosing Left
Value Difference (L - R)

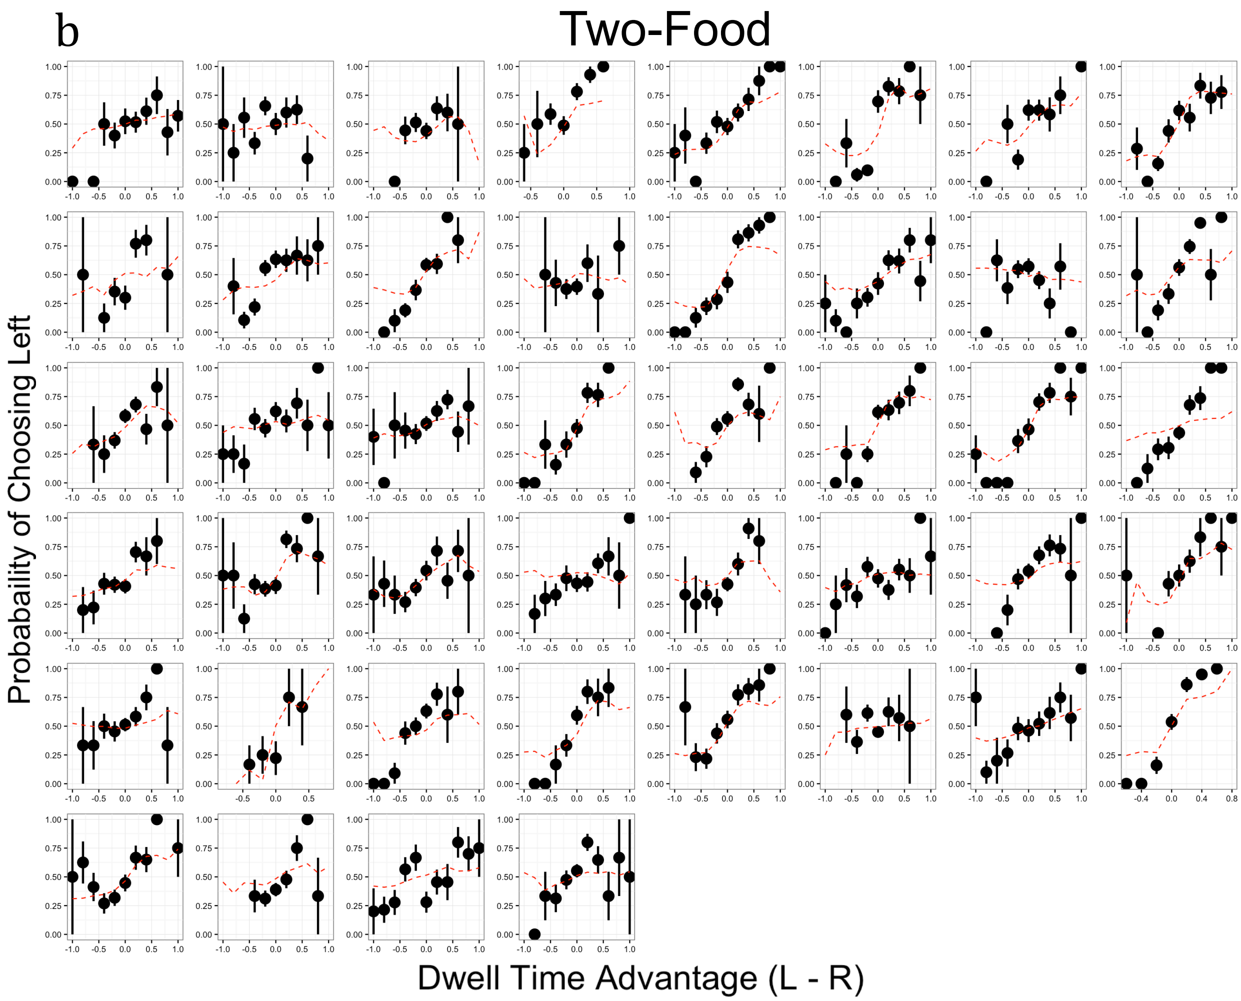
b
Two-Food
Probabaility of Choosing Left
Dwell Time Advantage (L - R)

c

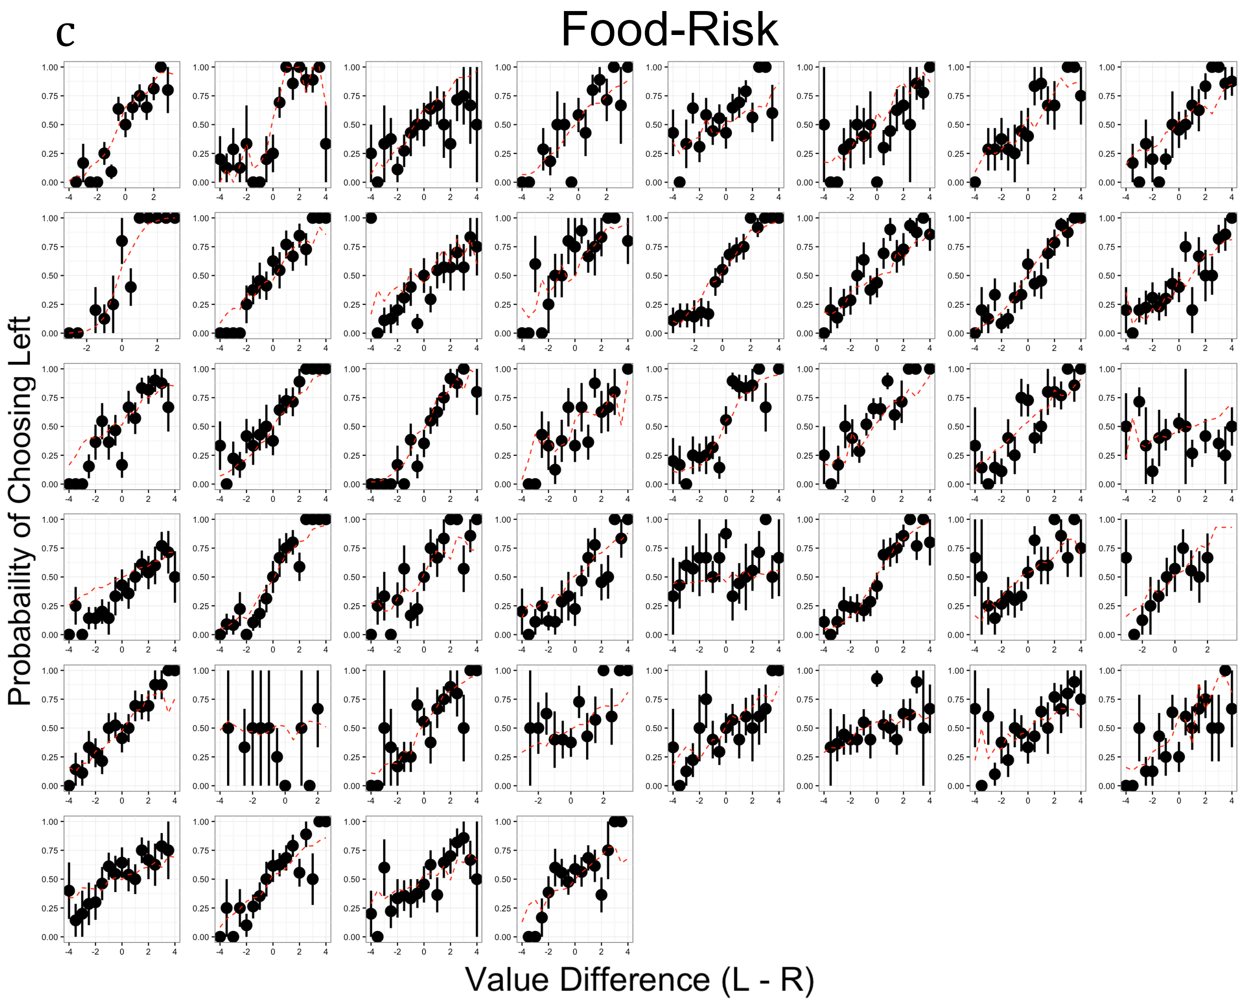

Food-Risk
Probabaility of Choosing Left
Value Difference (L - R)

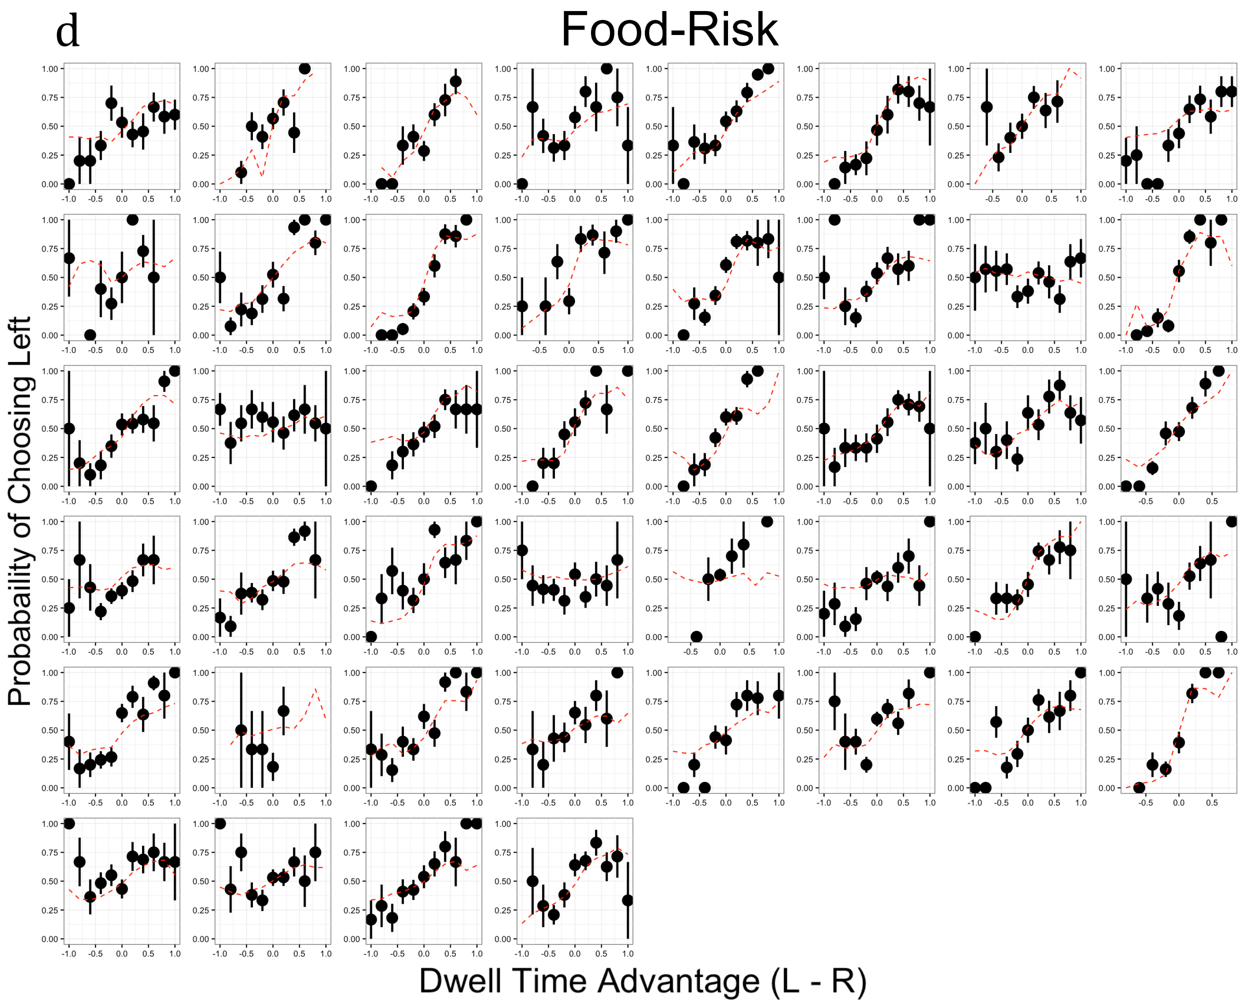
d
Food-Risk
Probabaility of Choosing Left
Dwell Time Advantage (L - R)

e

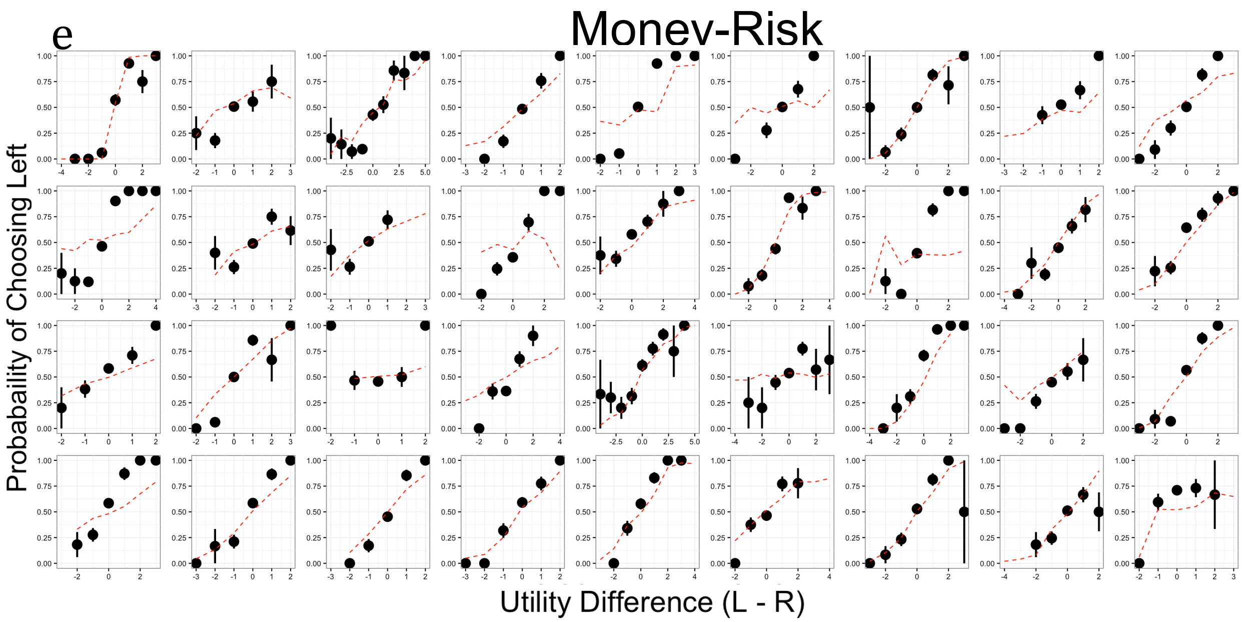
Money-Risk
Probabaility of Choosing Left
Utility Difference (L - R)

f

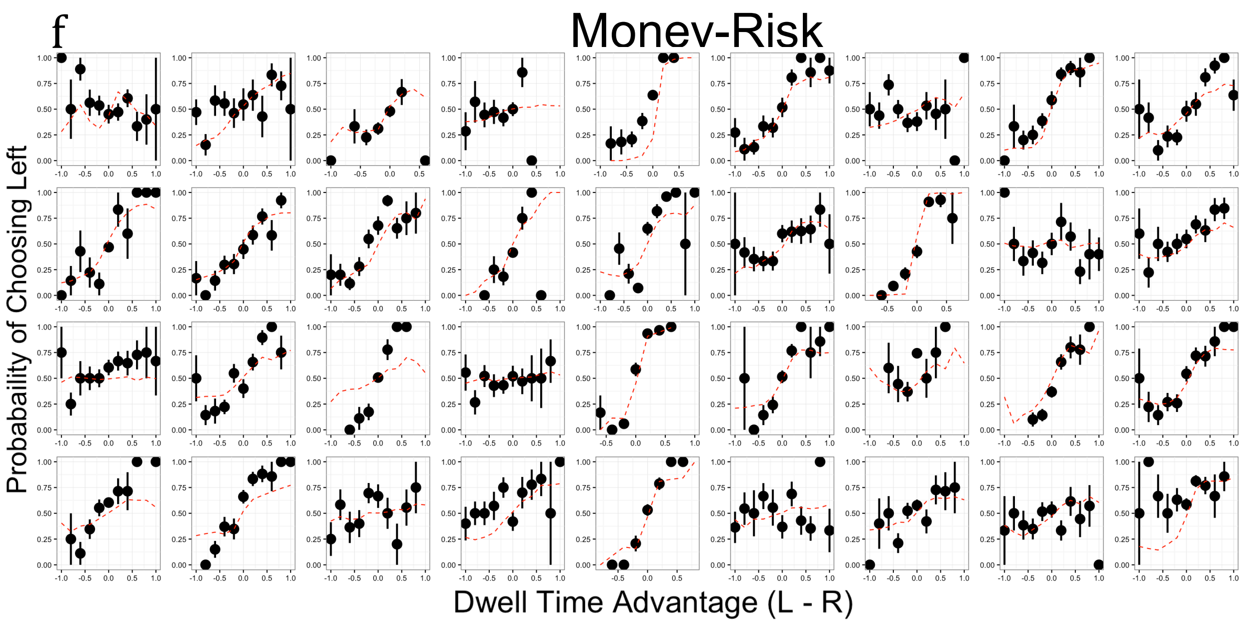
Money-Risk
Probabaility of Choosing Left
Dwell Time Advantage (L - R)

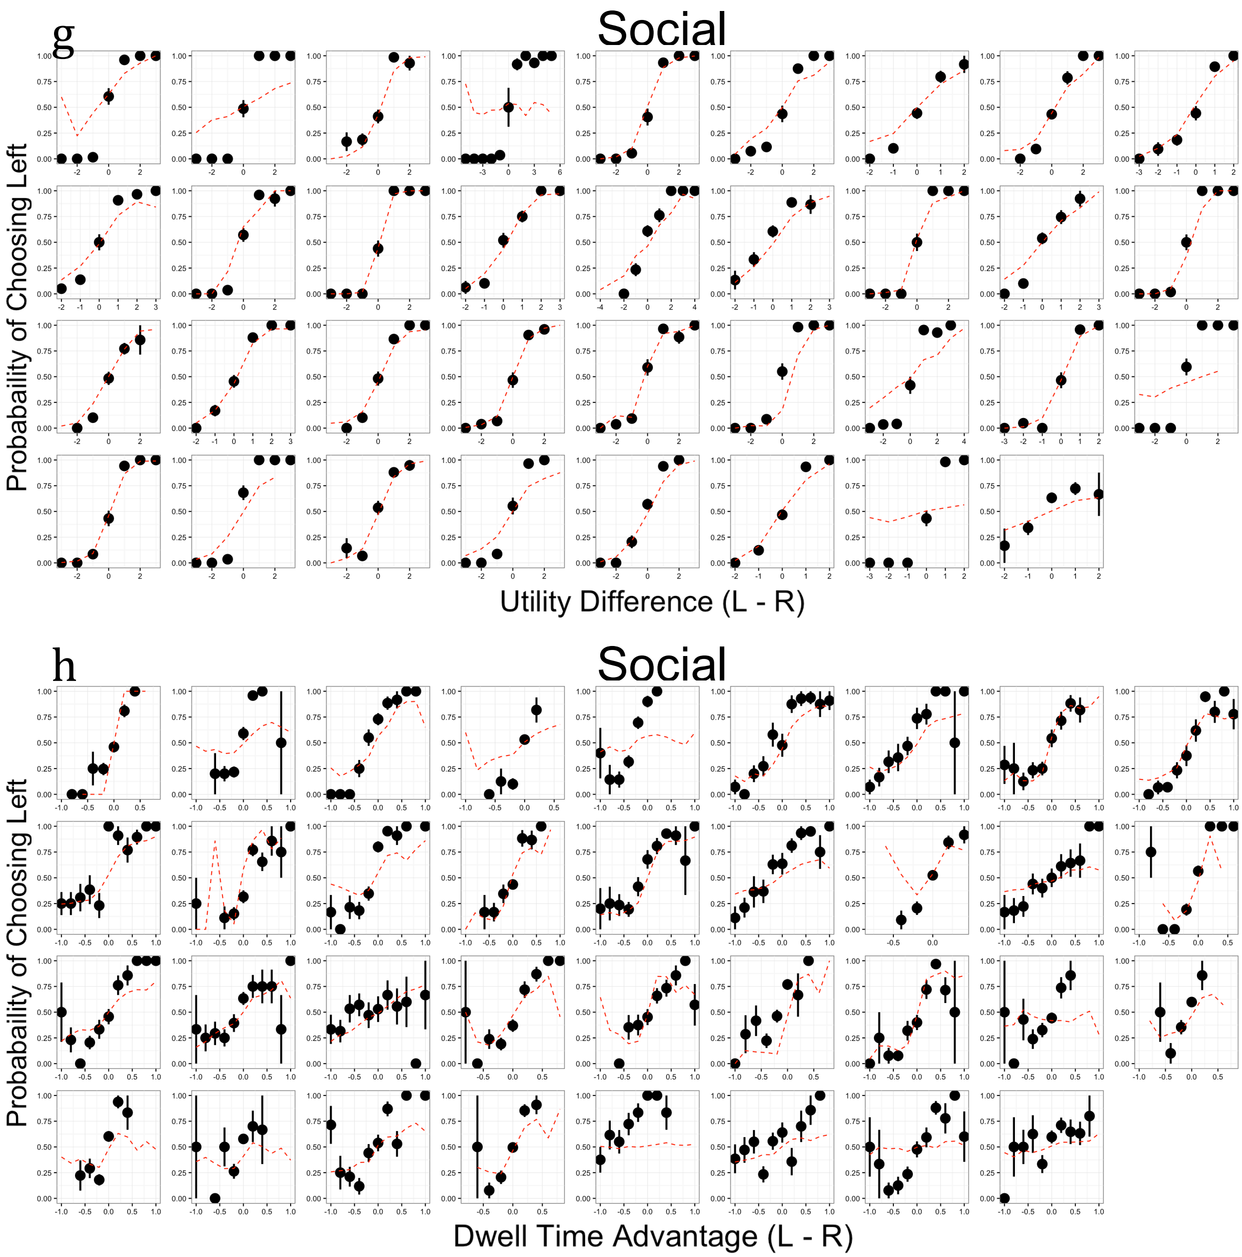

Figure S4. Individual aDDM fits. The individual-level fits for each task. Black circles represent data while red dashed lines represent aDDM simulations. Effect of utility on choice and effect of dwell time on choice for (a-b) two-food, (c-d) food-risk, (e-f) money-risk, and (g-h) social choice. Within each task, the individual fits are ordered, such that the first plot in the utility/choice mosaic and the first plot in the dwell time/choice mosaic come from the same subject. The dwell time plots demonstrate the necessity of controlling for utility when estimating logistic choice models; without it, the model can predict non-monotonic choice curves.

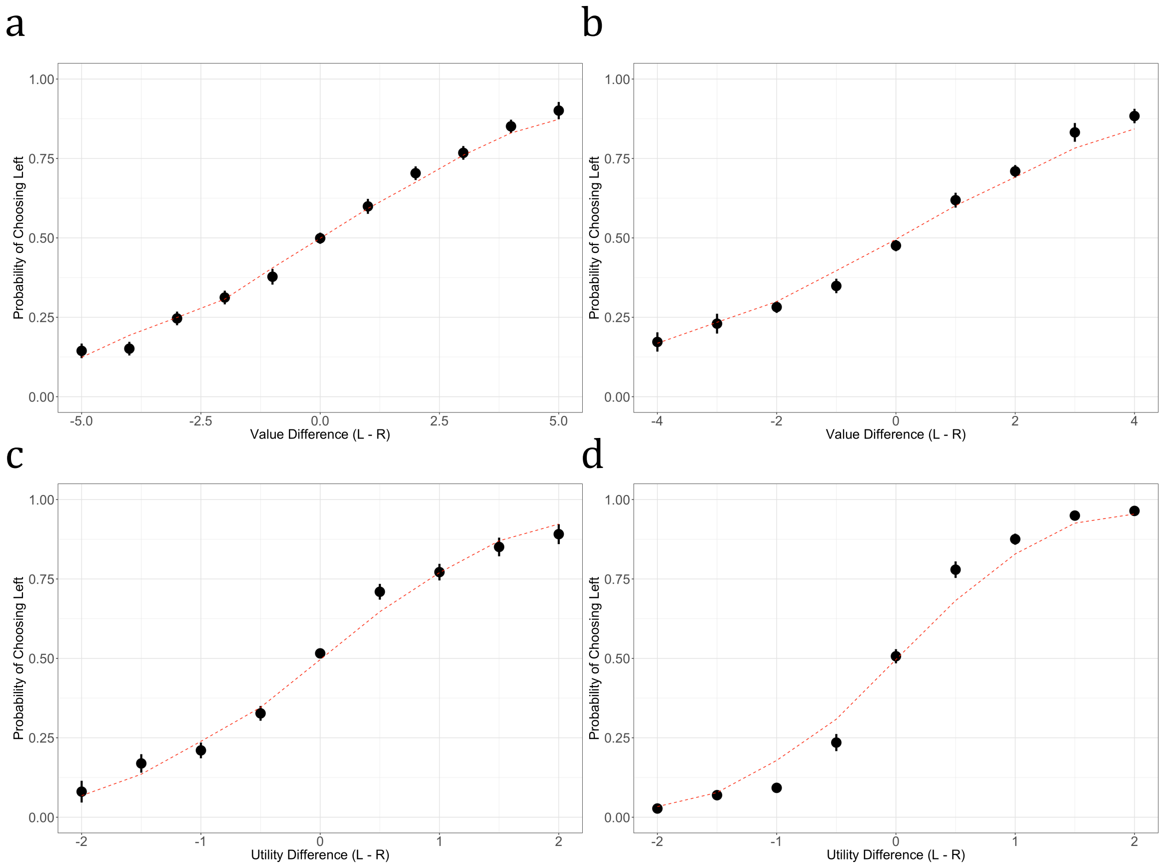

Figure S5. Choice-value consistency. Probability of choosing the left option vs. subjective value difference in the a) two-food, b) food-risk, c) money-risk, and d) social tasks. Black circles represent data while red dashed lines represent aDDM simulations. These simulations use randomly sampled dwell times and first dwell locations from the data. Each trial is simulated 10 times (with noise in the diffusion process) to create the simulated dataset.

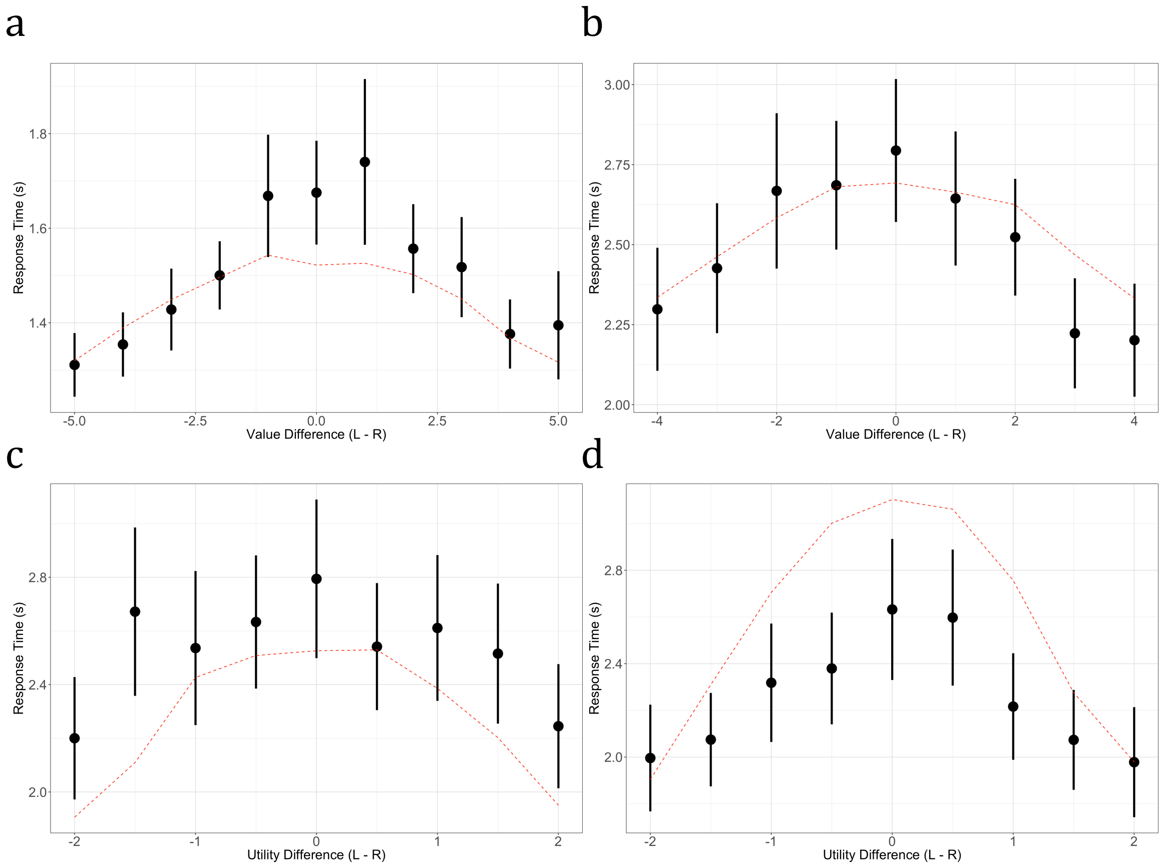

Figure S6. Response Times. RTs vs. subjective value difference in the a) two-food, b) food-risk, c) money-risk, and d) social tasks. Black circles represent data while red dashed lines represent aDDM simulations. These simulations use randomly sampled dwell times and first dwell locations from the data. Each trial is simulated 10 times (with noise in the diffusion process) to create the simulated dataset.

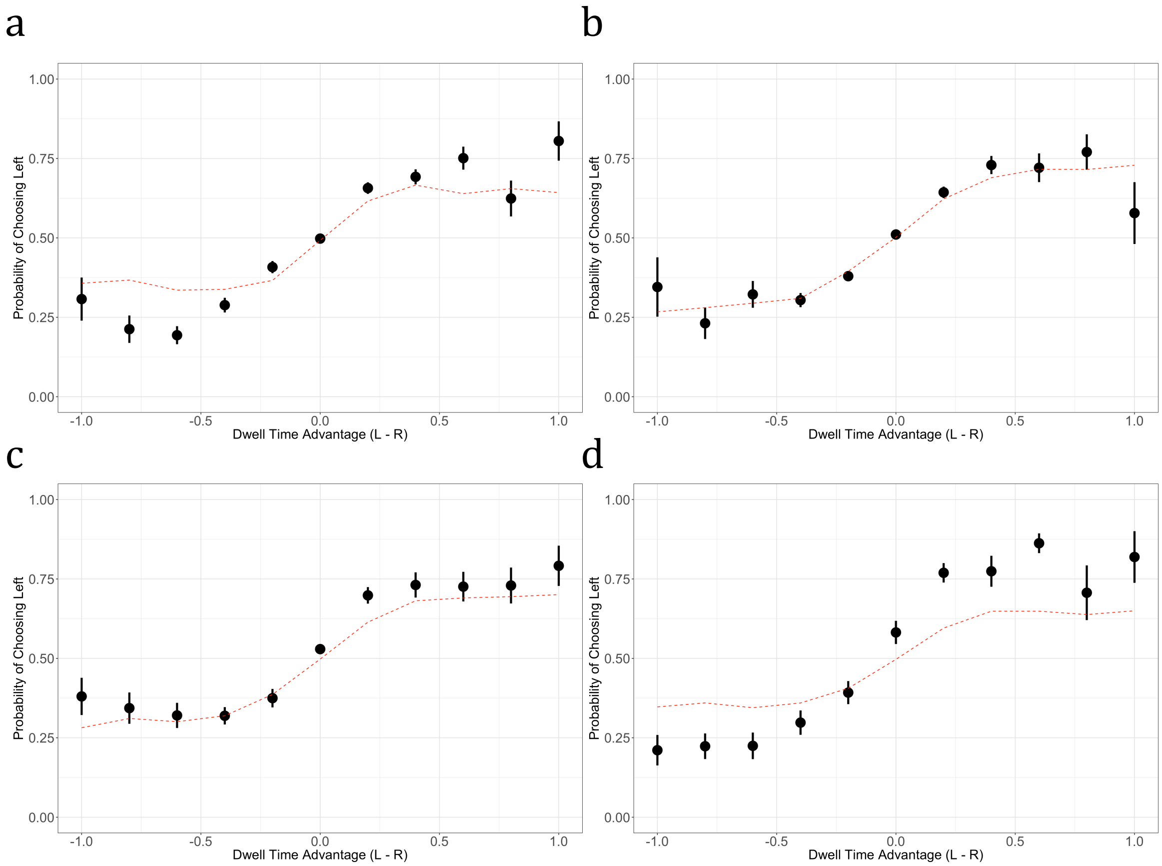

Figure S7. Effect of dwell time on choice. Choice vs. dwell time advantage in the a) two-food, b) food-risk, c) money-risk, and d) social tasks. Black circles represent data while red dashed lines represent aDDM simulations. These simulations use randomly sampled dwell times and first dwell locations from the data. Each trial is simulated 10 times (with noise in the diffusion process) to create the simulated dataset.

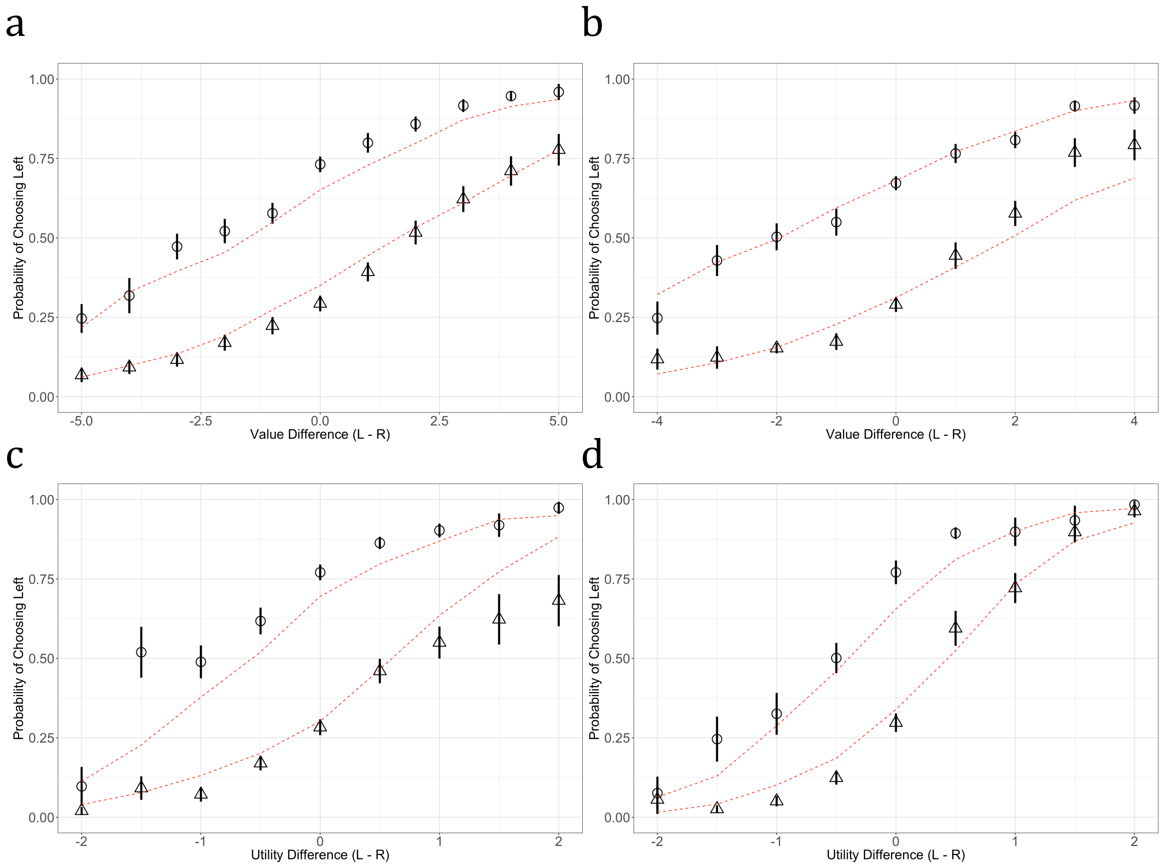

Figure S8. Final dwell effects. Choice vs. subjective value difference, conditioned on final dwell location in the a) two-food, b) food-risk, c) money-risk, and d) social tasks. Circles (triangles) represent final dwells to the left (right) option. Black circles and triangles represent data while red dashed lines represent aDDM simulations. These simulations use randomly sampled dwell times and first dwell locations from the data. Each trial is simulated 10 times (with noise in the diffusion process) to create the simulated dataset.

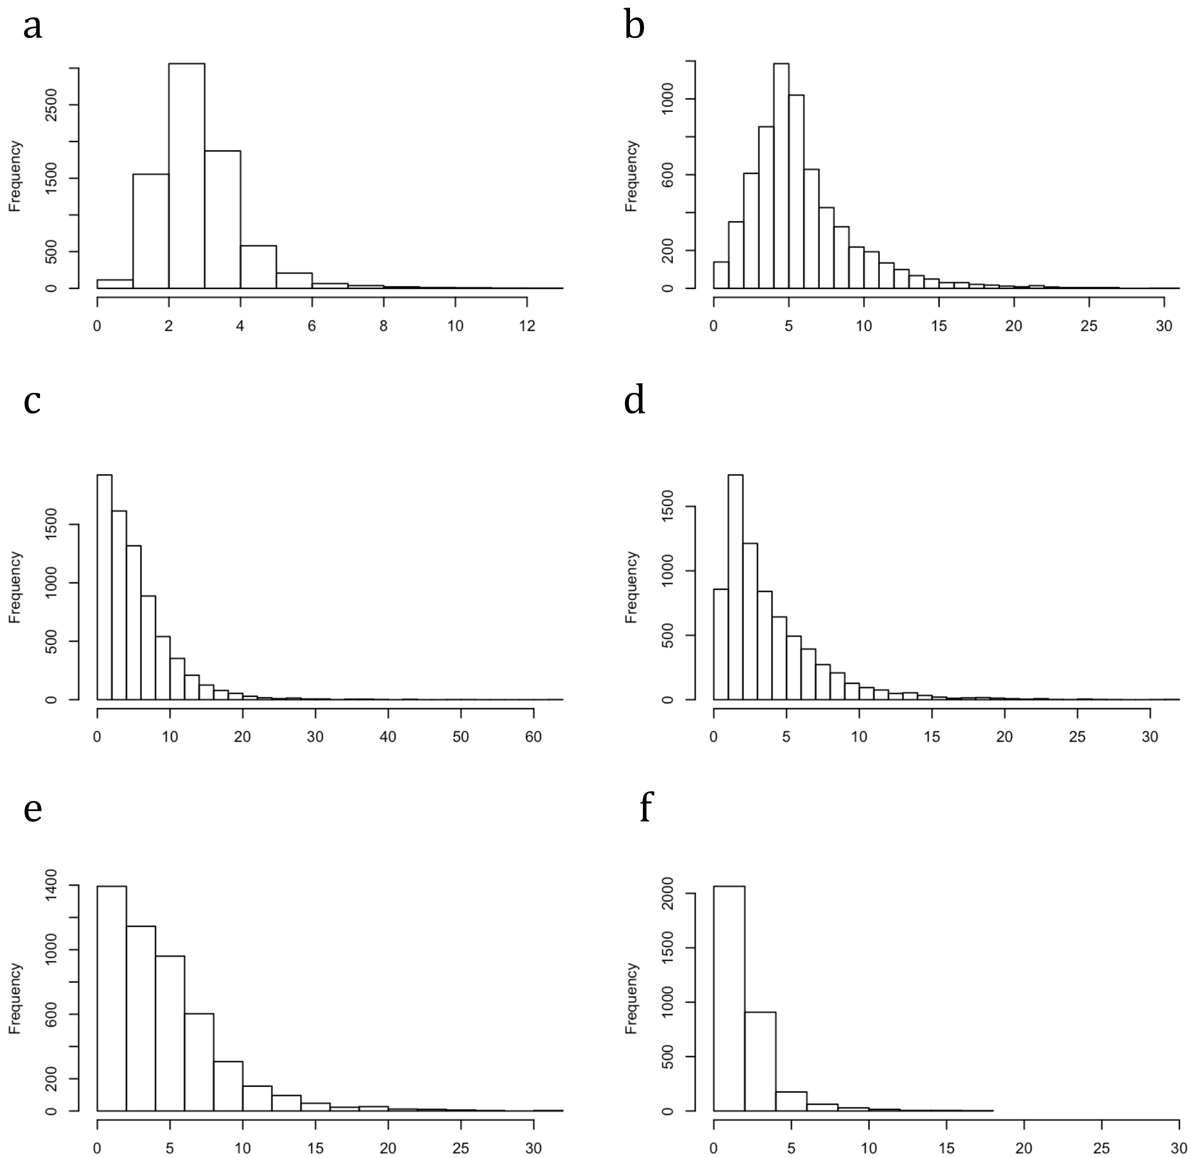

Figure S9. Number of dwells. Histograms of number of dwells in each trial in the a) two-food, b) food-risk, c) money-risk, d) social (combined), e) social (pro-social), and f) social (selfish) tasks. Every trial is included as one observation.

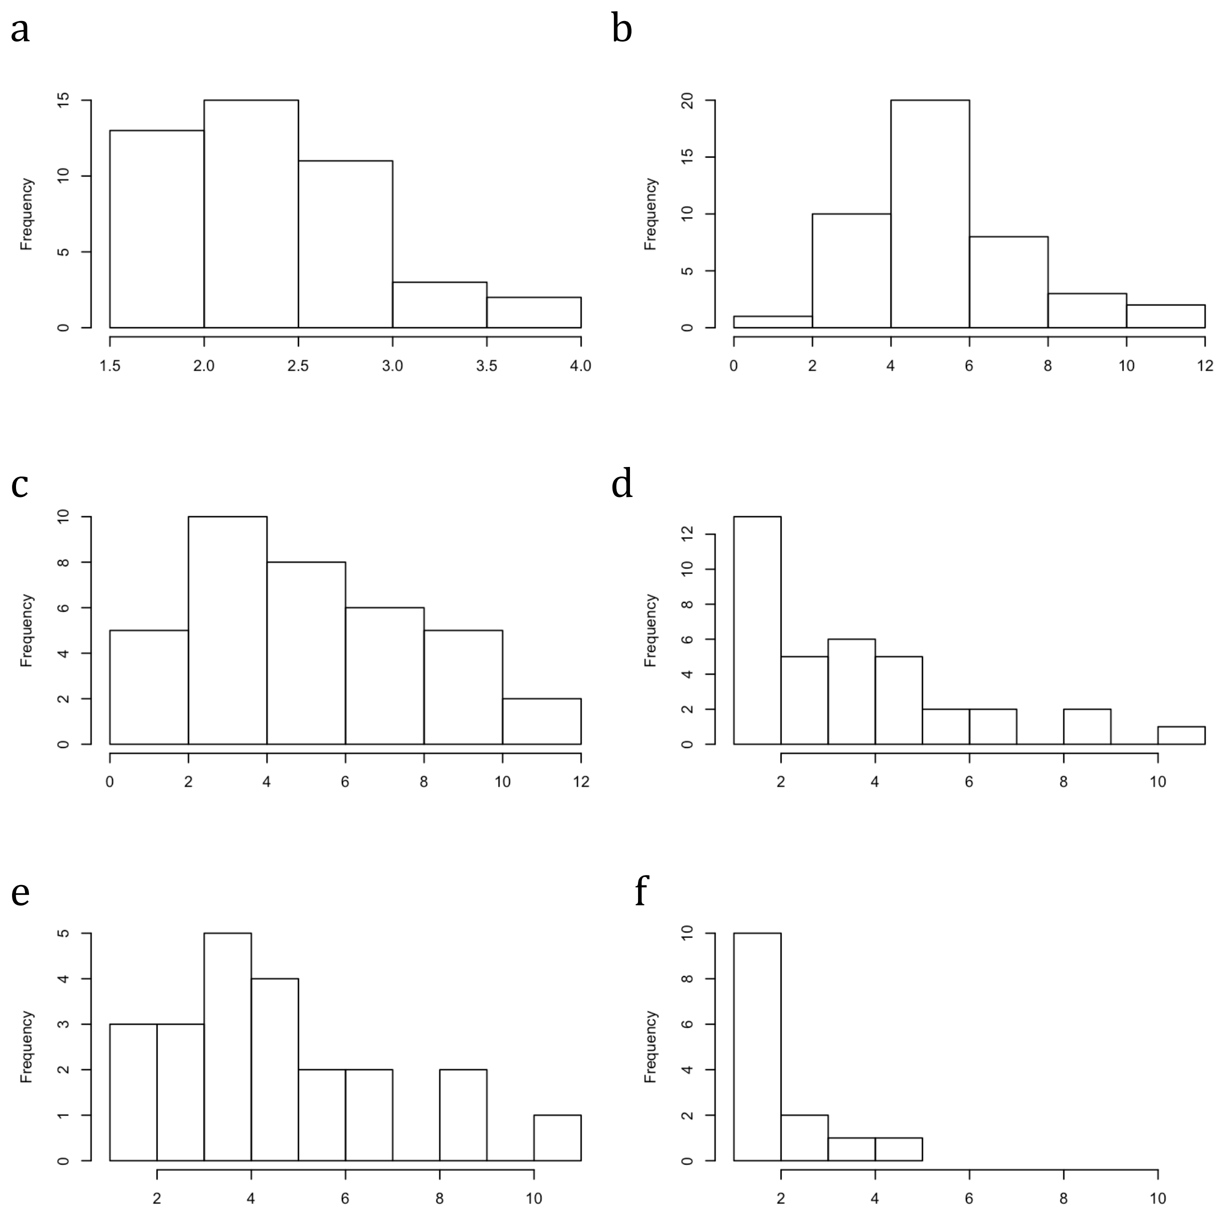

Figure S10. Average number of dwells. Histograms of average number of dwells per subject in the a) two-food, b) food-risk, c) money-risk, d) social (combined), e) social (pro-social), and f) social (selfish) tasks. Number of dwells are averaged at the subject level.

a b

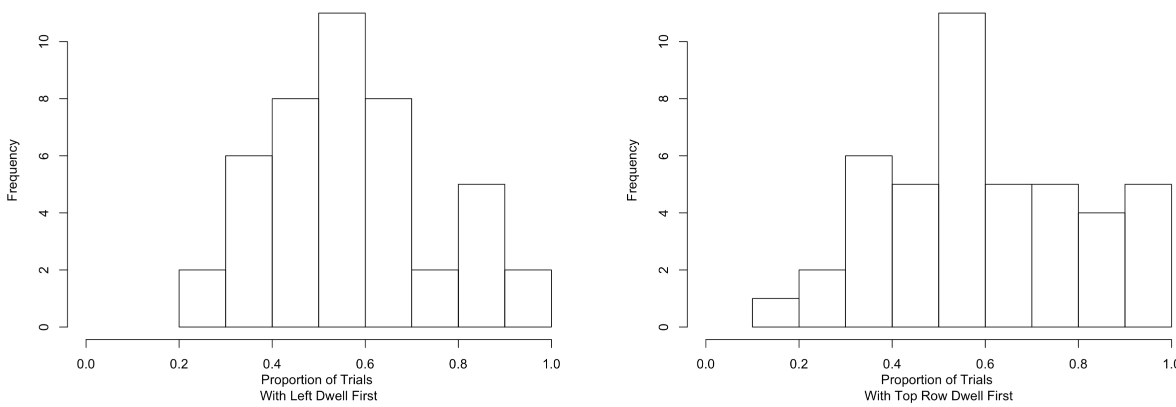

Figure S11. Subject-level differences in first dwell location. a) Subjects vary in their average tendency (across all four tasks) to look at the left option first. b) Subjects vary in their average tendency (across the three tasks with multiple rows) to look at the top row first.

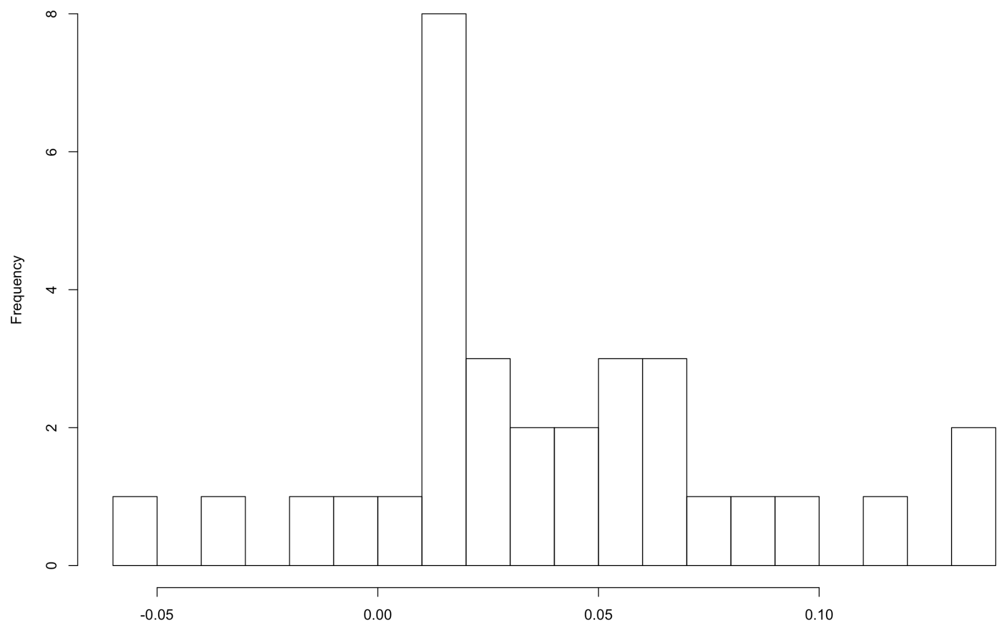

Figure S12. Histogram of sharpness gradients.

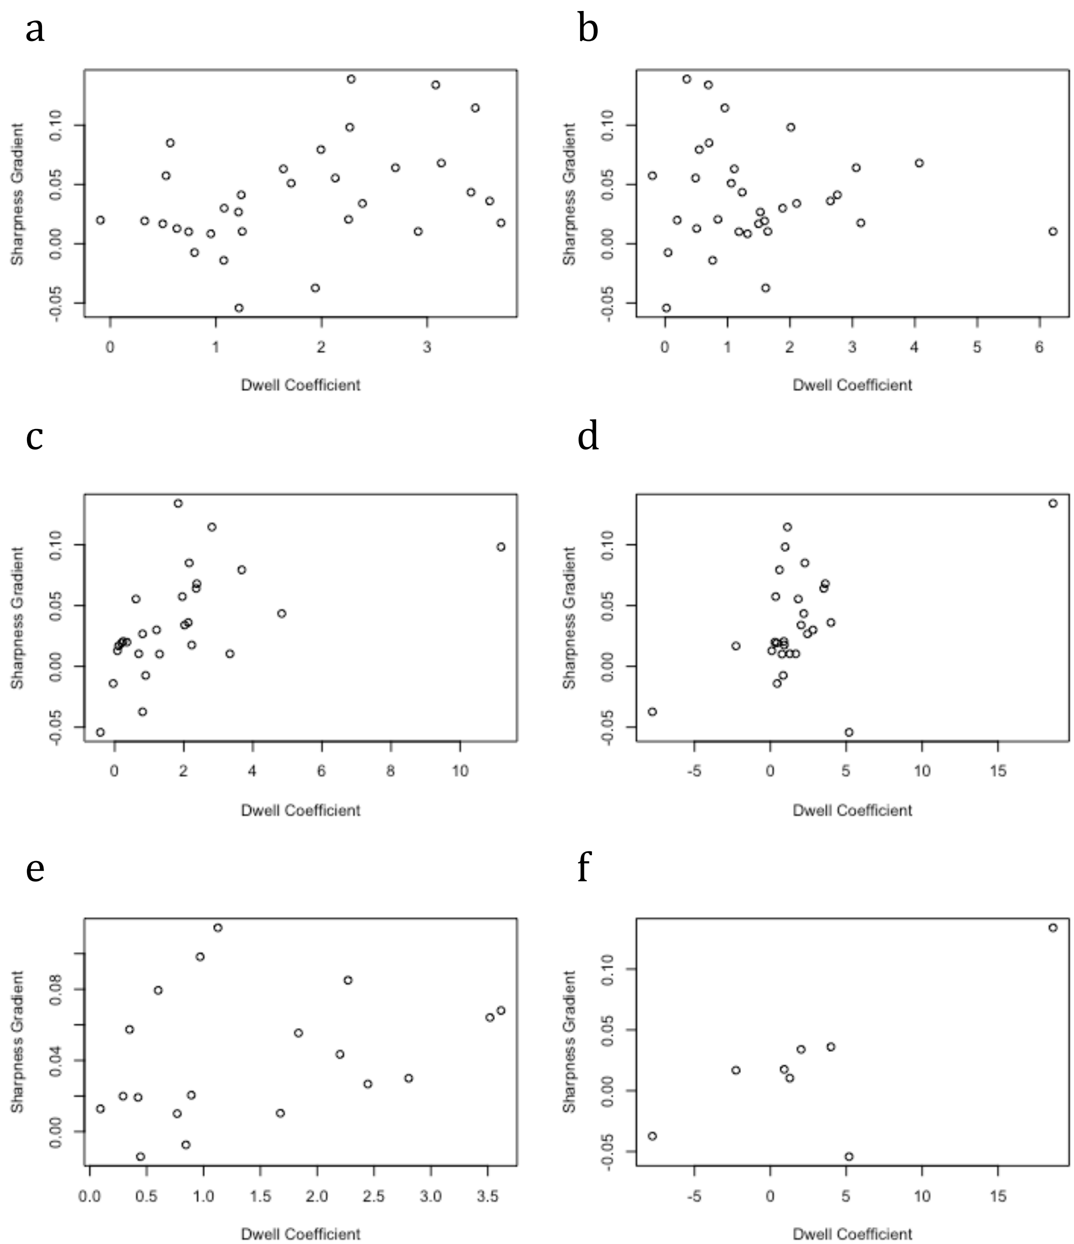

Figure S13. Sharpness and dwell time. Relationships between individual dwell time coefficients and sharpness gradients in the a) two-food, b) food-risk, c) money-risk, d) social (combined), e) social (pro-social), and f) social (selfish) tasks. The model regresses choosing the left option as a function of the value difference between the options and the dwell time advantage for the left option.

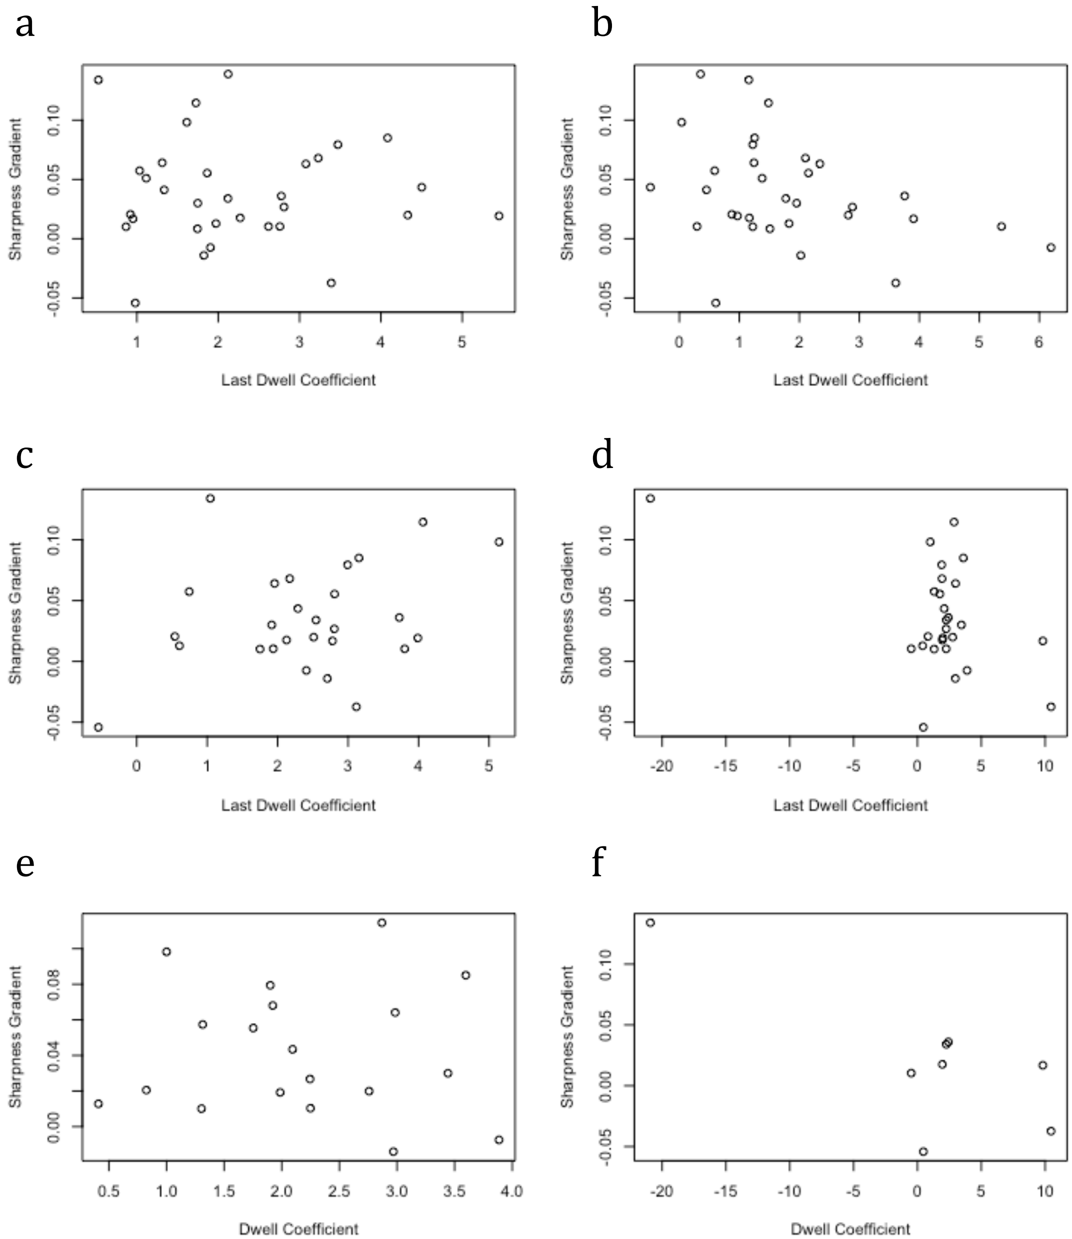

Figure S14. Sharpness and final dwell. Relationships between individual final dwell coefficients and sharpness gradients in the a) two-food, b) food-risk, c) money-risk, d) social (combined), e) social (pro-social), and f) social (selfish) tasks. The model regresses choosing the left option as a function of the value difference between the options and the location of the final dwell.

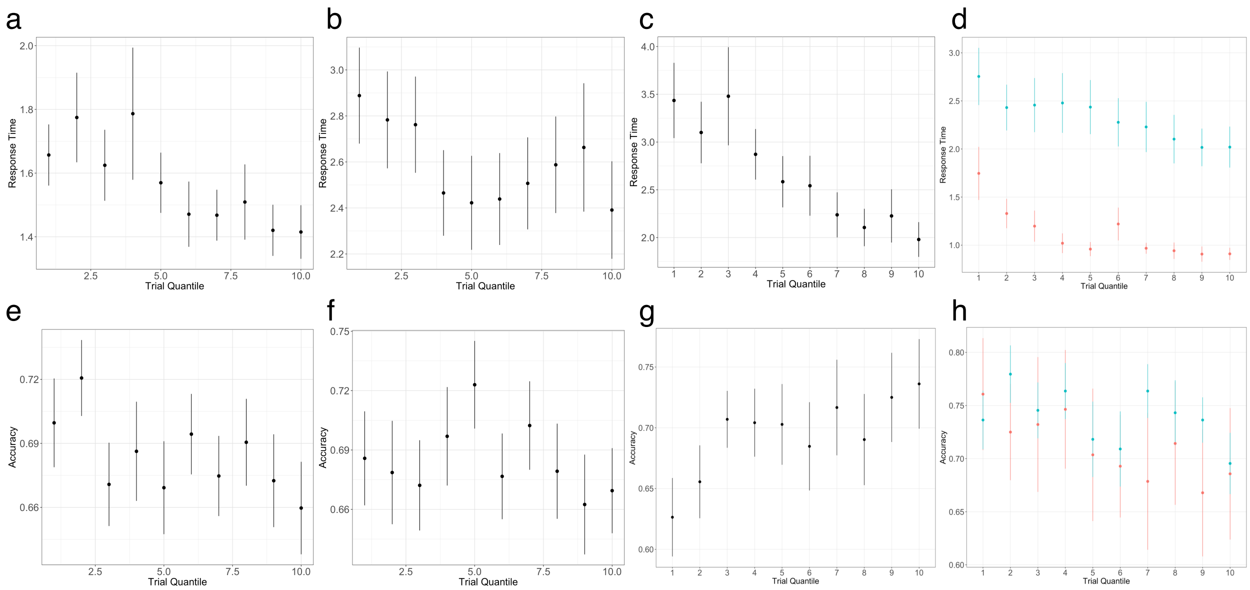

Figure S15. Trial number effects. Effects of trial number on RT in the a) two-food, b) food-risk, c) money-risk, and d) social tasks. Effects of trial number on accuracy in the e) two-food, f) food-risk, g) money-risk, and h) social tasks. Bars are standard errors across subjects. In the social task, the blue circles are the pro-social subjects and the orange circles are the selfish subjects.

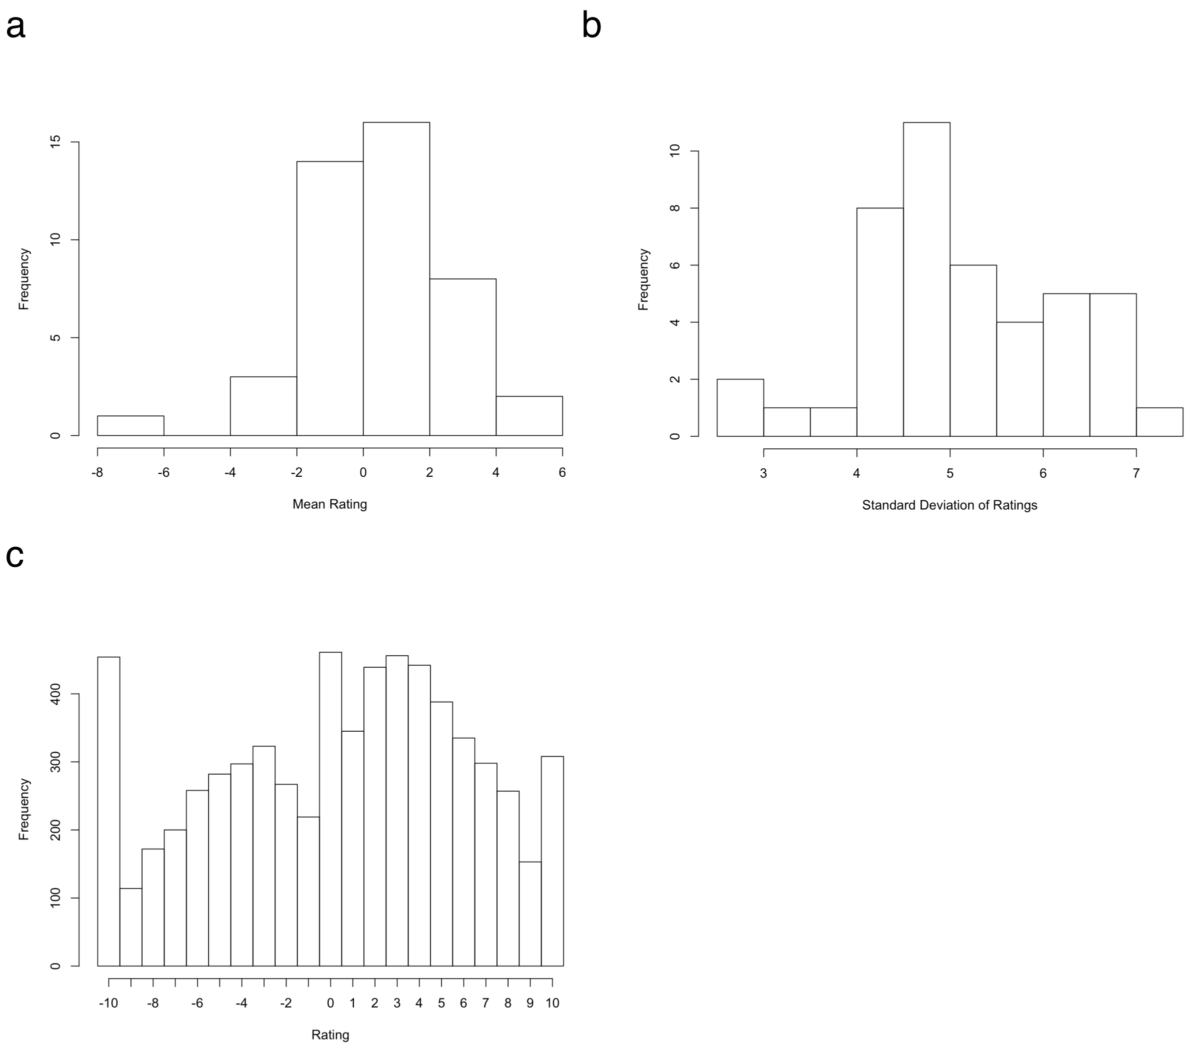

Figure S16. Rating Distributions. a) Histogram of mean rating by subject. Means are computed across all ratings at the subject level. b) Histogram of standard deviation of ratings by subject. Standard deviations are computed across all ratings at the subject level. c) Histogram of all ratings across all subjects.

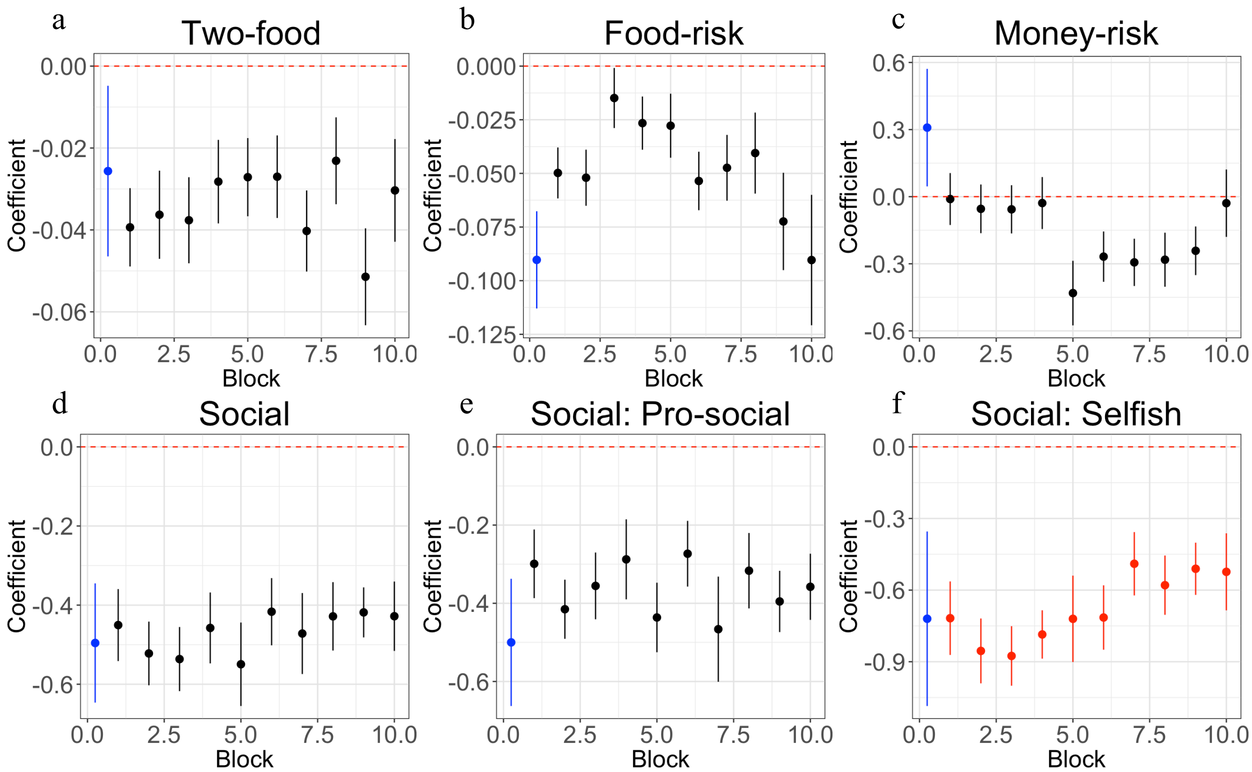

Figure S17. Coefficients on absolute value difference predicting logged RT, plotted as a function of trial block. Each block comprises 20 trials, except for the first (in blue) which just incudes the first 5 trials. Error bars represent standard errors on each coefficient. Toward the end of the tasks (i.e. later blocks), subjects' RTs are still inversely correlated with value difference in the a) two-food, b) food-risk, c) money-risk, d) social (combined), e) social (pro-social) and f) social (selfish) tasks (in red). However, the magnitude of the value difference/RT relationship decreases over time for the selfish subjects.

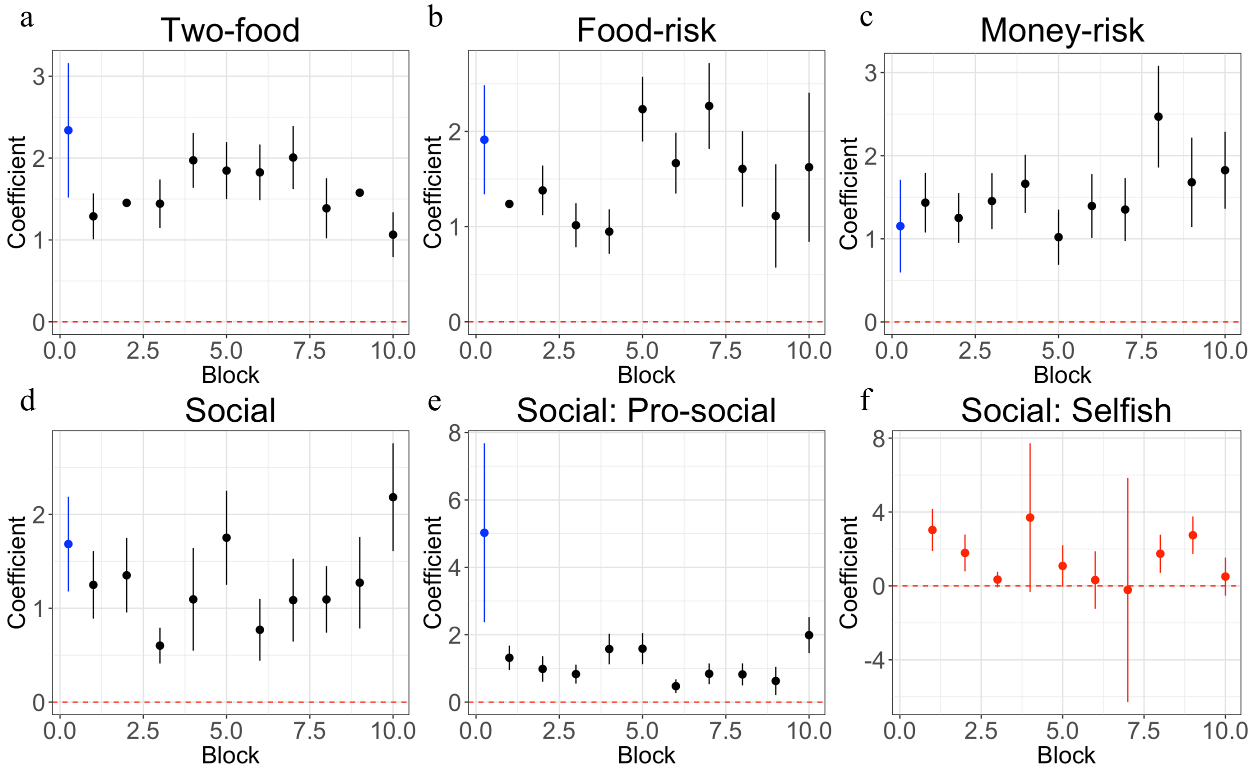

Figure S18. Coefficients on total dwell time advantage predicting choice, plotted as a function of trial block. Value difference is included as a covariate in the model: *Choice ~ Dwell Time Advantage + Value Difference.* Each block comprises 20 trials, except for the first (in blue) which just incudes the first 5 trials. Error bars represent standard errors on each coefficient. There do not seem to be any systematic time trends in the dwell-time coefficients in the a) two-food, b) food-risk, c) money-risk, d) social (combined) or e) social (pro-social) tasks, but there is a trend toward zero in the f) social (selfish) group. The estimate for the first five trials is missing from the selfish subjects as the standard error on the coefficient was ~30,000 and thus uninterpretable.

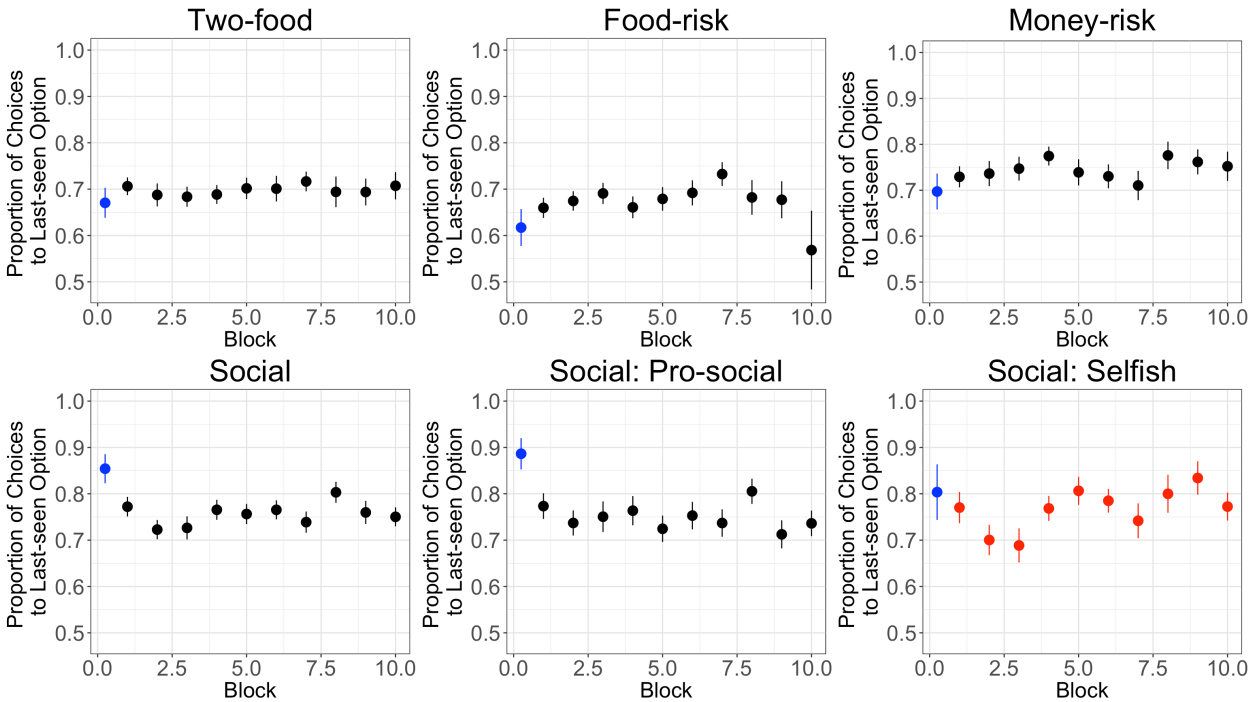

Figure S19. Proportion of choices to last-seen option as a function of trial number. Across the trials and across all tasks, there are not systematic changes in the proportion of last-gaze-to-chosen-option (all ps > 0.1), with the exception of the selfish subjects, who show a slight increase over time in the tendency to choose the last-seen option ($\beta = 0.041$, $p = 0.006$). Each block comprises 20 trials, except for the first (in blue) which just incudes the first 5 trials.

Table S1. Dwell time advantage coefficient correlations

Task	Two-Food	Food-Risk	Money-Risk
(All Participants, N = 36)			
Food-Risk	0.59***		
Money-Risk	0.54***	0.55***	
Social	0.46**	0.26	0.19
(Pro-social Participants Only, N = 22)			
Task	Two-Food	Food-Risk	Money-Risk
Food-Risk	0.58**		
Money-Risk	0.43*	0.43*	
Social	0.32	0.40’	0.08
(Selfish Participants Only, N = 14)			
Task	Two-Food	Food-Risk	Money-Risk
Food-Risk	0.66*		
Money-Risk	0.84***	0.75**	
Social	0.58*	0.09	0.38

‘$p < 0.10$ *$p < 0.05$ **$p < 0.01$ ***$p < 0.001$

Correlations are calculated within subjects across tasks. The model includes the signed value difference between the options, as well as the location of the final dwell (interacted with the absolute value difference).

Table S2. Final dwell coefficient correlations

(All Participants, N = 36)			
Task	Two-Food	Food-Risk	Money-Risk
Food-Risk	0.42*		
Money-Risk	0.42*	0.40*	
Social	0.18	0.40*	0.37*
(Pro-social Participants Only, N = 22)			
Task	Two-Food	Food-Risk	Money-Risk
Food-Risk	0.34		
Money-Risk	0.37’	0.37’	
Social	0.16	0.26	0.44*
(Selfish Participants Only, N = 14)			
Task	Two-Food	Food-Risk	Money-Risk
Food-Risk	0.53*		
Money-Risk	0.42	0.38	
Social	0.20	0.42	0.41
‘$p < 0.10$ *$p < 0.05$ **$p < 0.01$ ***$p < 0.001$			

Correlations are calculated within subjects across tasks. The model includes the signed value difference between the options, the total dwell time advantage, as well as the absolute value difference (interacted with the location of final dwell).

Table S3. Effects of task order on RT, accuracy, and utility parameters.

Task Location	Two-Food		Food-risk		Money-risk			
	RT	Accuracy	RT	Accuracy	RT	Accuracy	Rho	Risky Choice
1st	1.75 (0.29)	0.71 (0.03)	3.45 (0.41)	0.70 (0.01)	2.05 (0.29)	0.72 (0.03)	0.47 (0.15)	0.51 (0.09)
2nd	1.46 (0.24)	0.65 (0.01)	2.25 (0.20)	0.69 (0.02)	3.76 (0.81)	0.72 (0.05)	0.42 (0.16)	0.48 (0.10)
3rd	1.71 (0.14)	0.69 (0.02)	1.78 (0.23)	0.64 (0.04)	2.72 (0.54)	0.70 (0.03)	0.64 (0.19)	0.62 (0.07)
4th	1.37 (0.08)	0.68 (0.02)	2.40 (0.31)	0.69 (0.04)	2.36 (0.32)	0.71 (0.03)	0.58 (0.14)	0.58 (0.08)

Task Location	Social (pro-social)					Social (selfish)				
	RT	Accuracy	Alpha	Beta	Social Choice	RT	Accuracy	Alpha	Beta	Social Choice
1st	2.51 (0.79)	0.92 (0.02)	0.10 (0.05)	0.33 (0.02)	0.46 (0.05)	0.94 (0.04)	0.98 (0.01)	0.003 (0.003)	0.05 (0.01)	0.05 (0.02)
2nd	2.05 (0.29)	0.92 (0.01)	0.13 (0.07)	0.32 (0.03)	0.49 (0.09)	1.08 (0.17)	0.99 (0.01)	0.01 (0.01)	0.04 (0.03)	0.03 (0.02)
3rd	2.74 (0.93)	0.88 (0.02)	0.13 (0.06)	0.16 (0.10)	0.36 (0.05)	1.08 (0.14)	0.99 (0.003)	0.01 (0.01)	0.02 (0.01)	0.04 (0.01)
4th	2.40 (0.51)	0.90 (0.02)	0.03 (0.02)	0.27 (0.04)	0.33 (0.04)	1.46 (0.36)	0.98 (0.01)	0.03 (0.003)	0.01 (0.01)	0.03 (0.02)

Table S4. Regression of percentage of choices for last-seen option as a function of trial.

Task	Coefficient (SE)
Two-food	0.0063 (0.0098)
Food-risk	0.013 (0.012)
Money-risk	0.0090 (0.0089)
Social	0.013 (0.010)
Social: Selfish	0.041** (0.015)
Social: Pro-Social	–0.0037 (0.013)
** $p < 0.01$. All other ps > 0.1.	